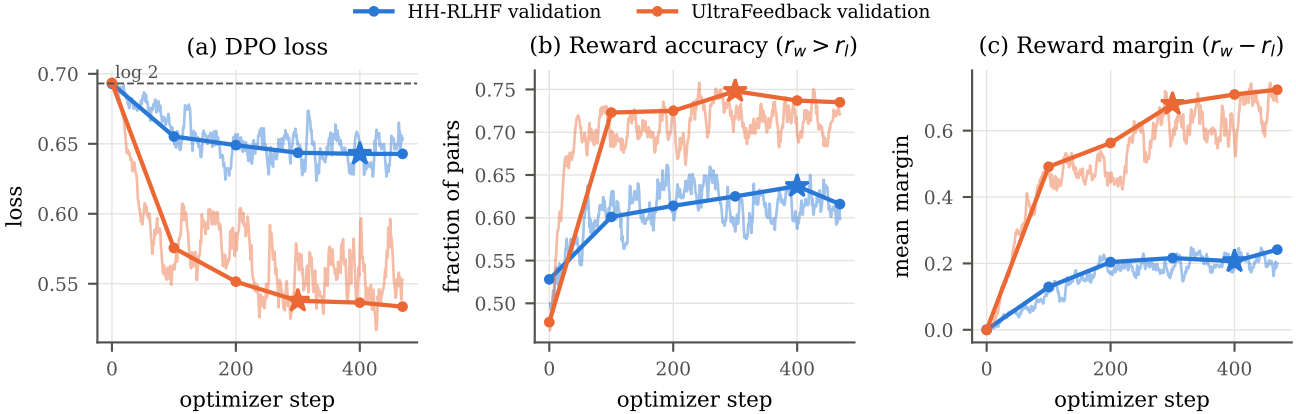
HH-RLHF validation
UltraFeedback validation
(a) DPO loss
log 2
loss
optimizer step
(b) Reward accuracy ($r_w > r_l$)
fraction of pairs
optimizer step
(c) Reward margin ($r_w - r_l$)
mean margin
optimizer step