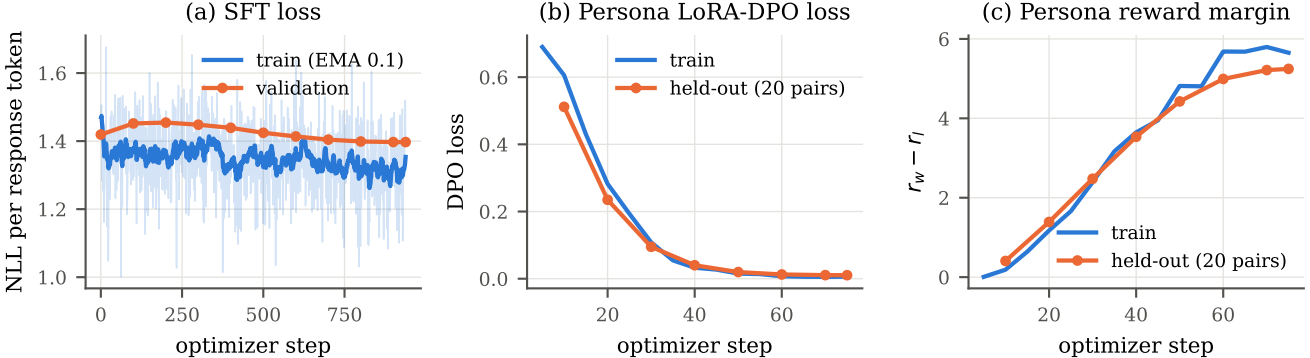
(a) SFT loss
NLL per response token
1.0
1.2
1.4
1.6
0
250
500
750
optimizer step
train (EMA 0.1)
validation
(b) Persona LoRA-DPO loss
DPO loss
0.0
0.2
0.4
0.6
20
40
60
optimizer step
train
held-out (20 pairs)
(c) Persona reward margin
$r_w - r_l$
0
2
4
6
20
40
60
optimizer step
train
held-out (20 pairs)